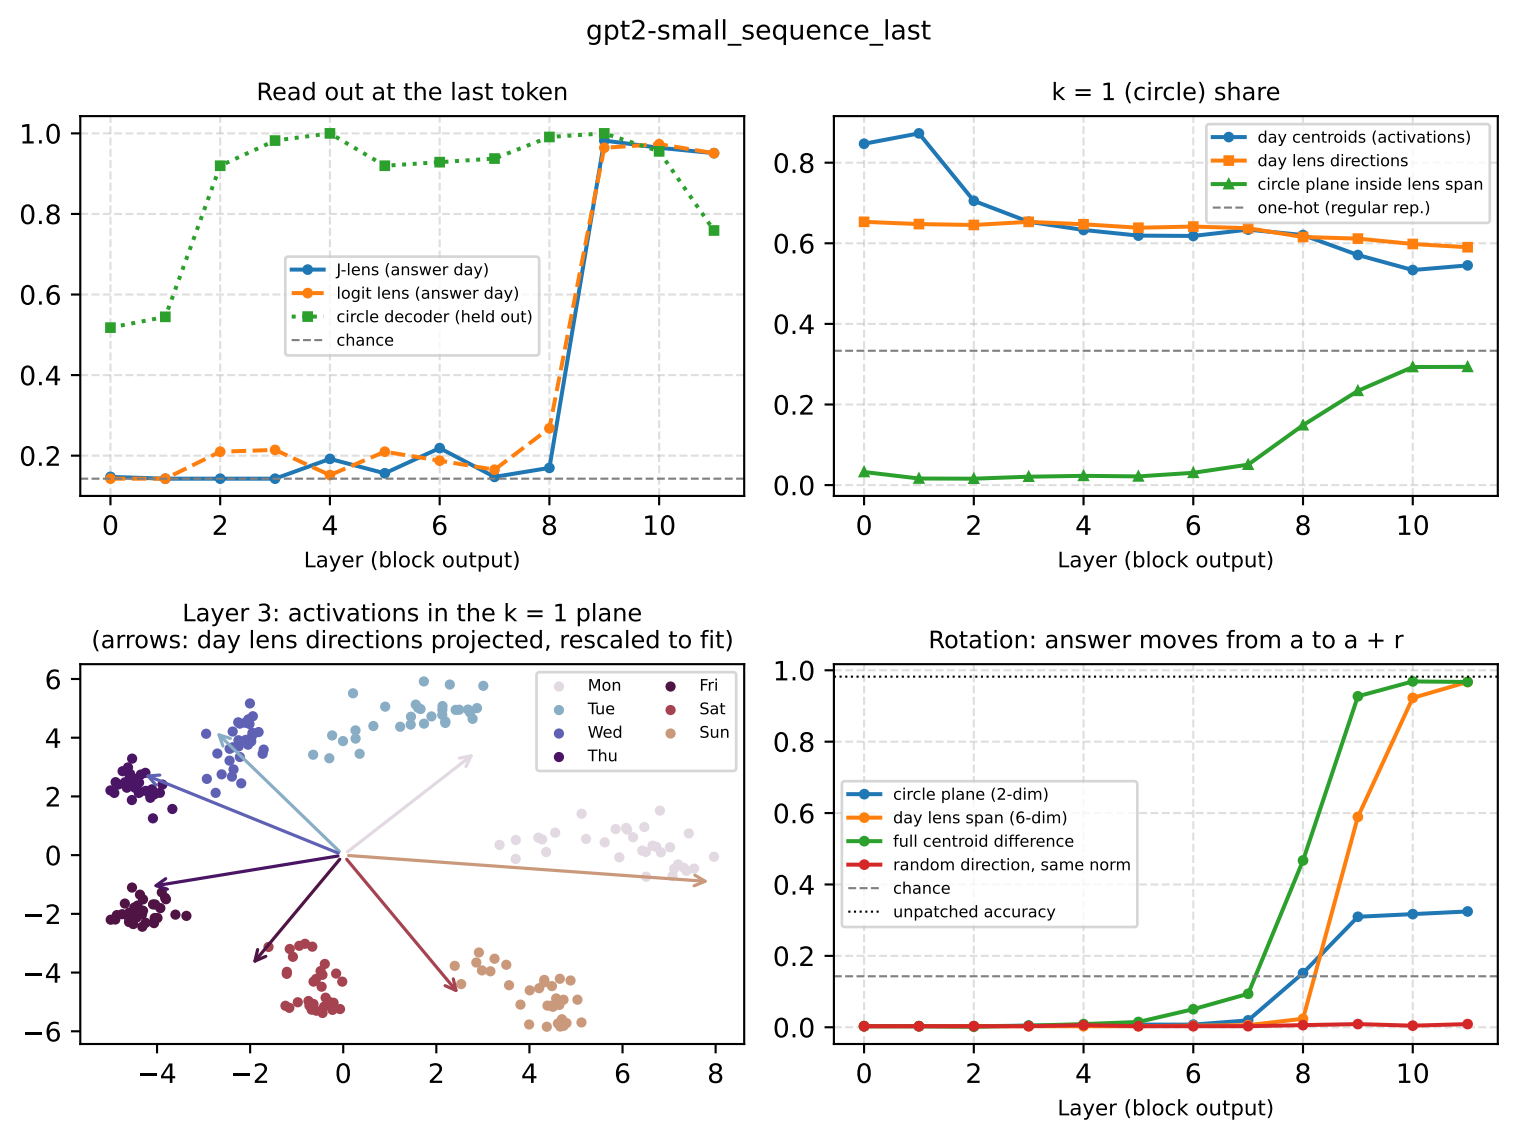
gpt2-small_sequence_last
Read out at the last token
J-lens (answer day)
logit lens (answer day)
circle decoder (held out)
chance
Layer (block output)
k = 1 (circle) share
day centroids (activations)
day lens directions
circle plane inside lens span
one-hot (regular rep.)
Layer (block output)
Layer 3: activations in the k = 1 plane
(arrows: day lens directions projected, rescaled to fit)
Mon
Tue
Wed
Thu
Fri
Sat
Sun
Rotation: answer moves from a to a + r
circle plane (2-dim)
day lens span (6-dim)
full centroid difference
random direction, same norm
chance
unpatched accuracy
Layer (block output)