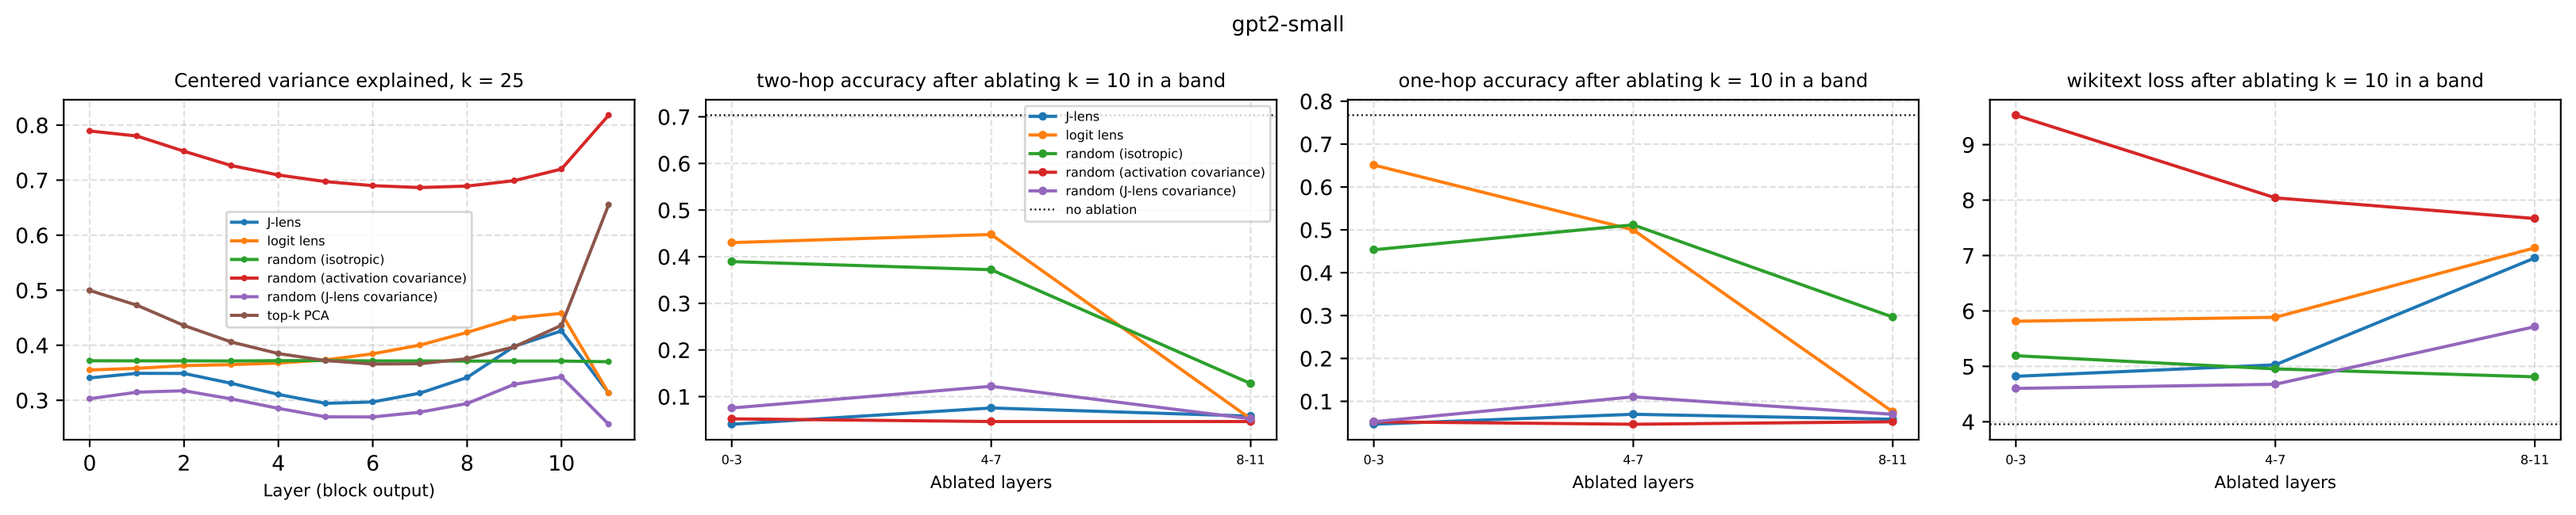
gpt2-small
Centered variance explained, k = 25
J-lens
logit lens
random (isotropic)
random (activation covariance)
random (J-lens covariance)
top-k PCA
Layer (block output)
two-hop accuracy after ablating k = 10 in a band
no ablation
Ablated layers
one-hop accuracy after ablating k = 10 in a band
Ablated layers
wikitext loss after ablating k = 10 in a band
Ablated layers
0-3
4-7
8-11