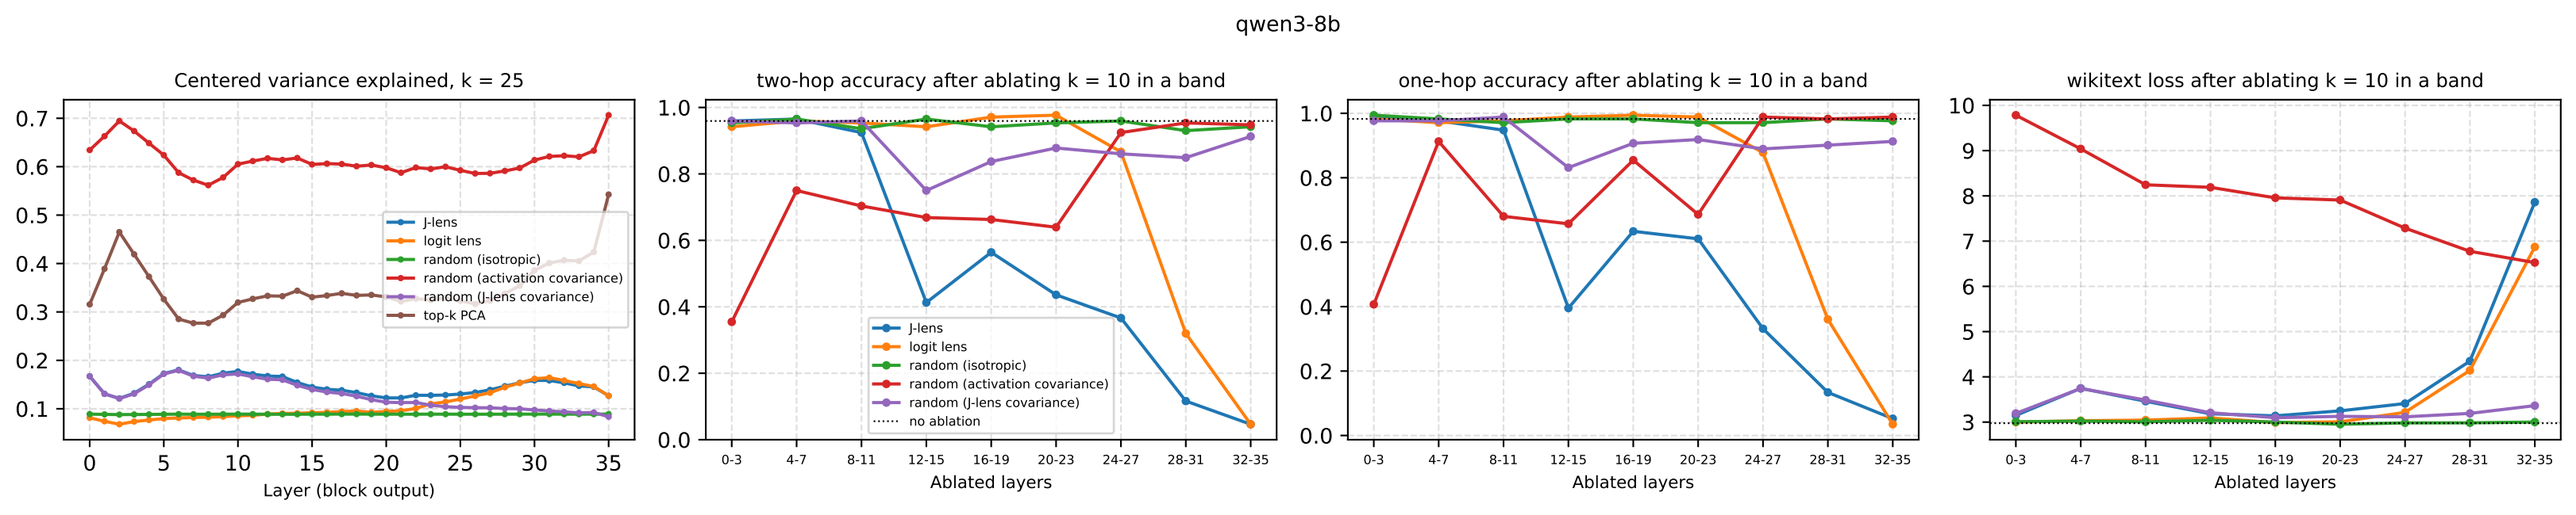
qwen3-8b
Centered variance explained, k = 25
Layer (block output)
J-lens
logit lens
random (isotropic)
random (activation covariance)
random (J-lens covariance)
top-k PCA
two-hop accuracy after ablating k = 10 in a band
Ablated layers
no ablation
one-hop accuracy after ablating k = 10 in a band
wikitext loss after ablating k = 10 in a band
0-3
4-7
8-11
12-15
16-19
20-23
24-27
28-31
32-35